

# **Class\_07\_lab**

Tessa Sterns PID: 18482353

An introductory exploration of machine learning. Starting with clustering and dimensionality reduction.

## **k-means clustering**

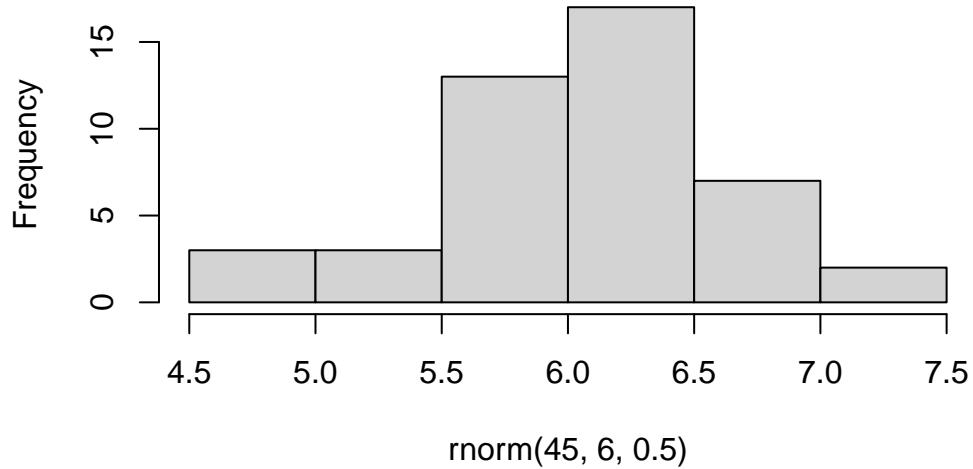
Using generated data from the `rnorm()` function with known clustering we can explore how the clustering works.

```
rnorm(20, 5, 2)
```

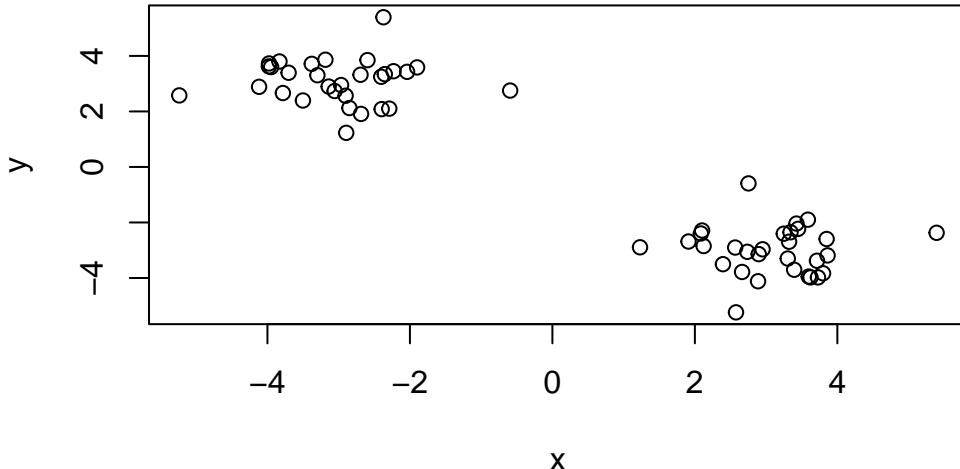
```
[1] 3.1996141 -0.1559747 5.7758461 3.4720581 2.9708914 6.0925479  
[7] 5.7590508 6.4455249 2.7798891 7.6198831 4.0639618 1.3986097  
[13] 6.2101924 5.0941210 3.8802428 2.7350357 5.4285824 1.8351155  
[19] 3.7168233 1.4435921
```

```
hist(rnorm(45, 6, 0.5))
```

## Histogram of rnorm(45, 6, 0.5)



```
x <- c(rnorm(30, -3),  
       rnorm(30, 3))  
  
y <- rev(x)  
  
z <- cbind(x,y)  
  
plot(z)
```



In base R the function `kmeans()` can do k-means clustering.

```
km <- kmeans(z, 2)
```

To retrieve the results from a returned list object use the \$ syntax

Q1 How many points are in each cluster?

km\$size

[1] 30 30

Q2 What is the cluster assignment?

km\$cluster

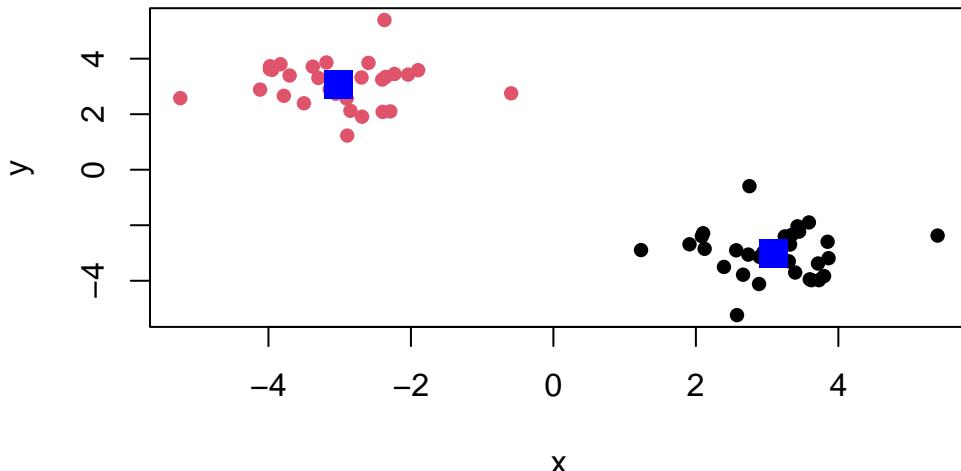
Q3 Where are the cluster centers?

```
km$centers
```

```
x           y  
1 3.083788 -3.009904  
2 -3.009904  3.083788
```

Q4 Make a clustering results figure of the data colored by cluster membership and show cluster centers?

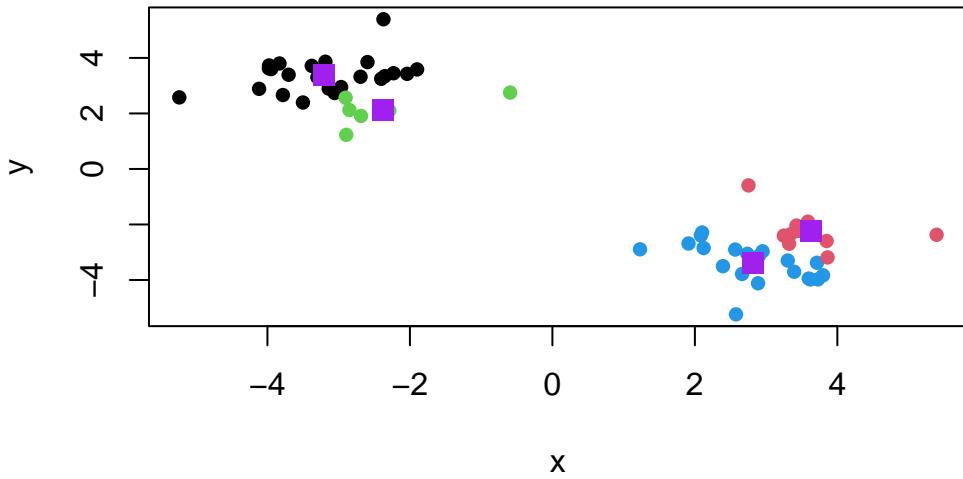
```
plot(z, col = km$cluster, pch = 16)  
points(km$centers, col="blue", pch = 15, cex = 2)
```



k-means clustering is popular, very fast and strait forward to use, taking numeric data as input and returns cluster membership vector. However you need to already know how many clusters exist in the data set.

Q5 Run k-means again with four groups and plot results?

```
km4 <- kmeans(z, 4)  
plot(z, col = km4$cluster, pch = 16)  
points(km4$centers, col="purple", pch = 15, cex = 1.5)
```



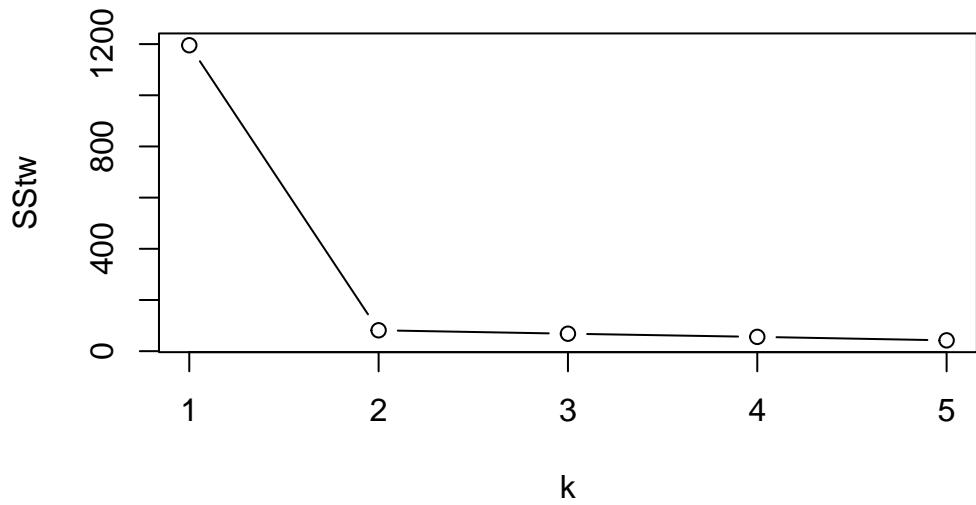
generating a Scree plot brute-force method

```

km3 <- kmeans(z, 3)
km5 <- kmeans(z, 5)
km1 <- kmeans(z, 1)

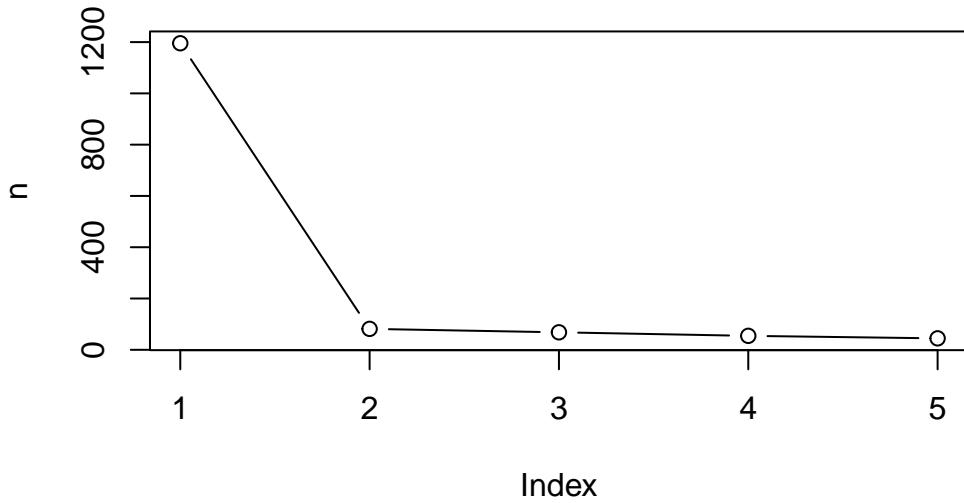
SStw <- c(km1$tot.withinss, km$tot.withinss, km3$tot.withinss, km4$tot.withinss, km5$tot.withinss)
k <- c(1,2,3,4,5)
plot(k, SStw, type = "b")

```



for loop to generate a Scree plot, more elegant code for when brute force is unreasonable.

```
n <- NULL
for(i in 1:5) {
  n <- c(n, kmeans(z, centers = i)$tot.withinss)
}
plot(n, type = "b")
```



## Hierarchical Clustering

The main “base” function is `hclust()`. Here we can’t imput raw data, need to generate a distance matrix using `dist()` function.

```
d <- dist(z)
hc <- hclust(d)
hc
```

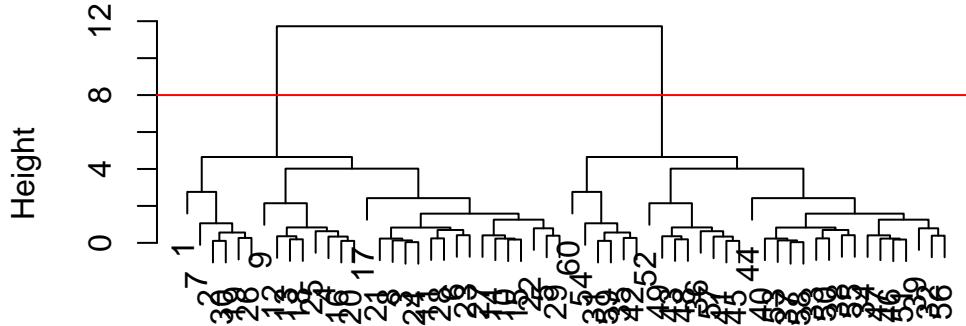
```
Call:
hclust(d = d)

Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

There is a plot method for `hclust` which will return a dendrogram.

```
plot(hc)
abline(h=8, col = "red")
```

## Cluster Dendrogram



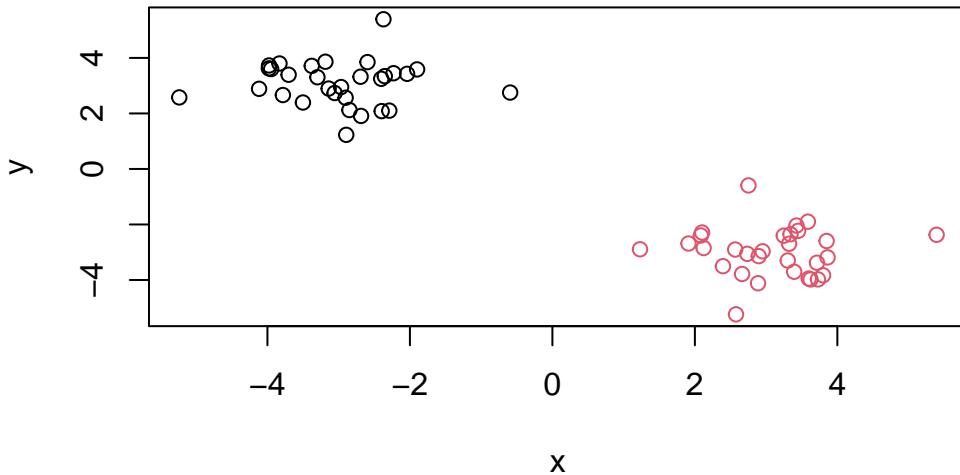
```
d  
hclust (*, "complete")
```

To retrieve our cluster membership vector we can use `cutree` at a given height 'h =' or groups 'k ='.

```
grps <- cutree(hc, h = 8)
```

Q6 Plot the data with our hclust result coloring.

```
plot(z, col = grps)
```



## Dimensionality reduction, (PCA) Principal Component Analysis

Principle components are new axis that are closest to the observations. They can capture more of the spread of the data. Useful for identifying outliers and trends in the data.

## Part 2 PCA

### PCA of UK food data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1 How many rows and columns are in this data set?

```
dim(x)
```

```
[1] 17 5
```

```
head(x)
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139

```
rownames(x) <- x[,1]  
x <- x[,-1]  
head(x)
```

		England	Wales	Scotland	N.Ireland
Cheese		105	103	103	66
Carcass_meat		245	227	242	267
Other_meat		685	803	750	586
Fish		147	160	122	93
Fats_and_oils		193	235	184	209
Sugars		156	175	147	139

```
# an easier way is to read in the row names in the first place, x <- read.csv(url, row.names=1)  
head(x)
```

		England	Wales	Scotland	N.Ireland
Cheese		105	103	103	66
Carcass_meat		245	227	242	267
Other_meat		685	803	750	586
Fish		147	160	122	93
Fats_and_oils		193	235	184	209
Sugars		156	175	147	139

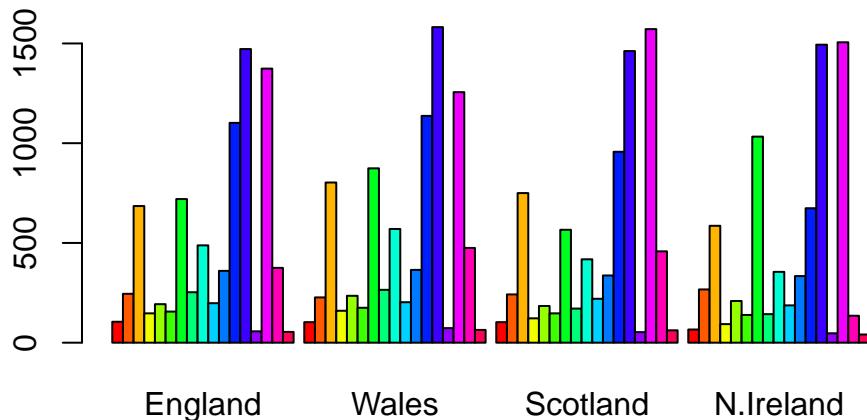
```
dim(x)
```

```
[1] 17 4
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

reading in the code properly the first time makes it so we don't accidentally remove extra columns.

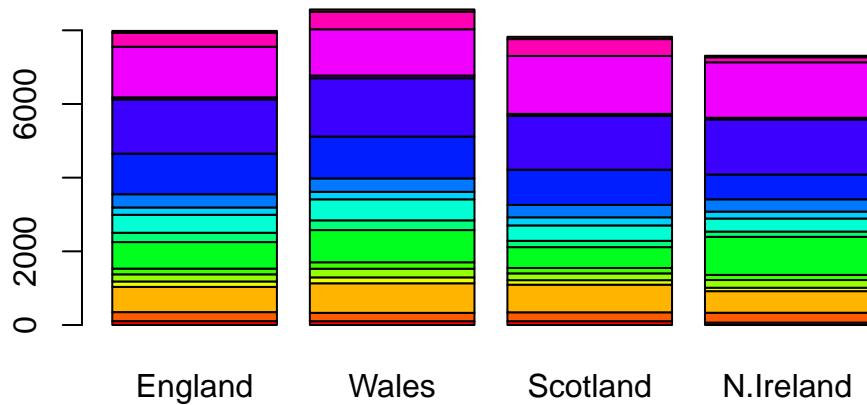
```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

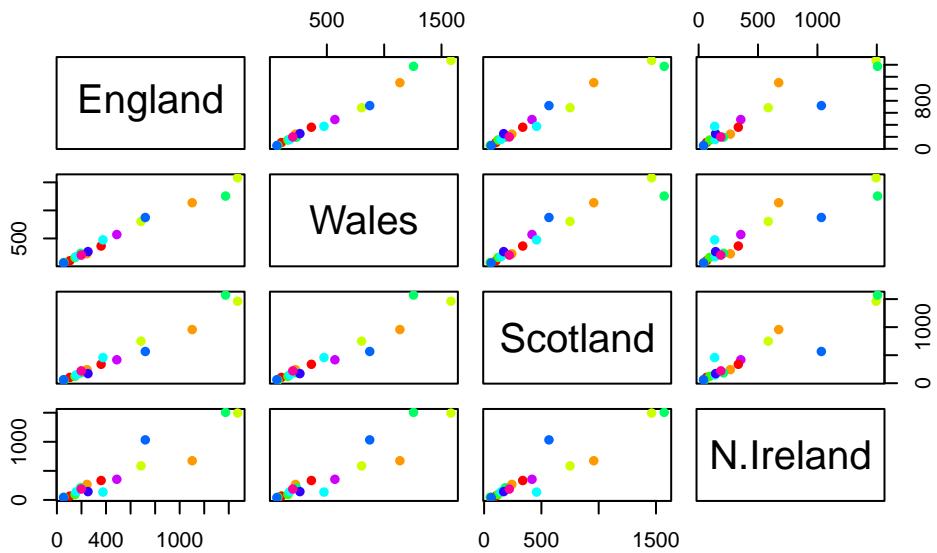
beside = F

```
barplot(as.matrix(x), col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



This code can generate all columns against each other. If they are on the diagonal then the countries eat the same amount of a given food.

Key takeaway: It is rather difficult to spot the major trends and patterns even in relatively small datasets. This would be absolutely useless with large datasets.

## PCA to save the day

The main function in “base” R for PCA is `prcomp()`. Transpose `t()` of the data set, rows become columns and columns become rows.

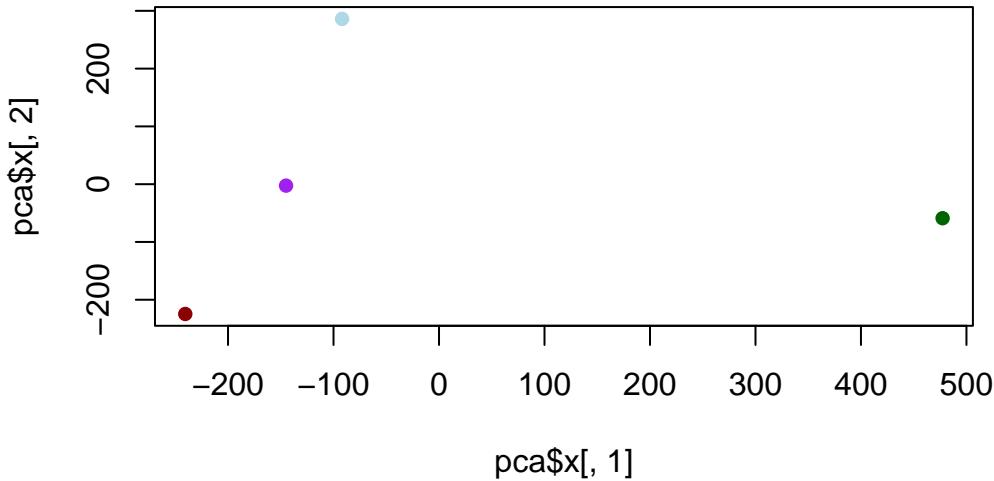
```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Base R plot

```
cols <- c("purple", "darkred", "lightblue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col = cols, pch =16)
```

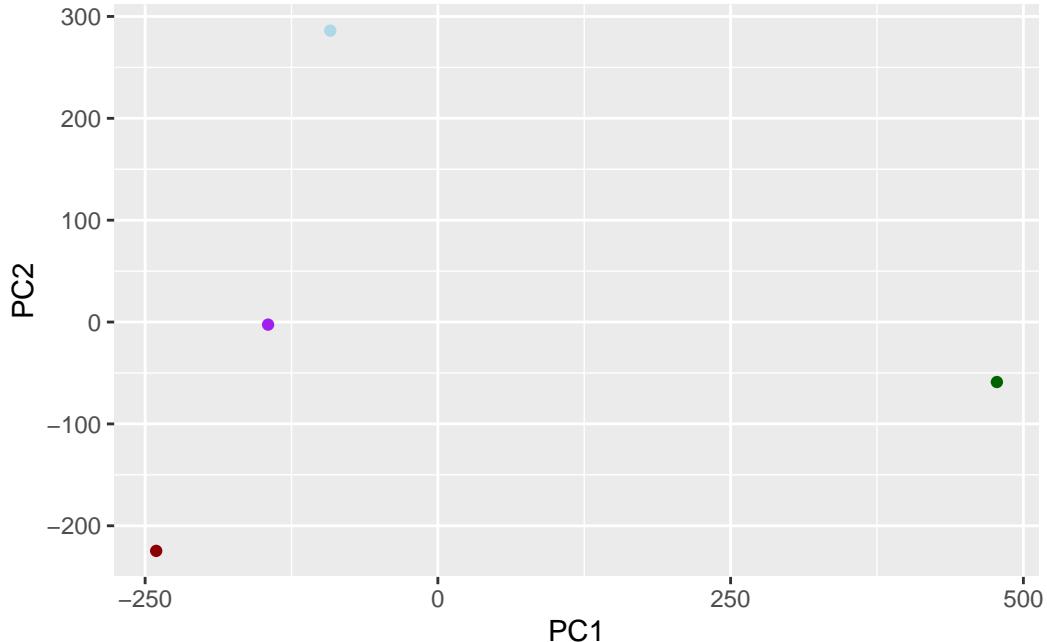


This shows that N.Ireland is different than the other groups.

However we can make a better graph

```
library(ggplot2)
```

```
ggplot(pca$x) +
  aes(PC1, PC2) +
  geom_point(col = cols)
```

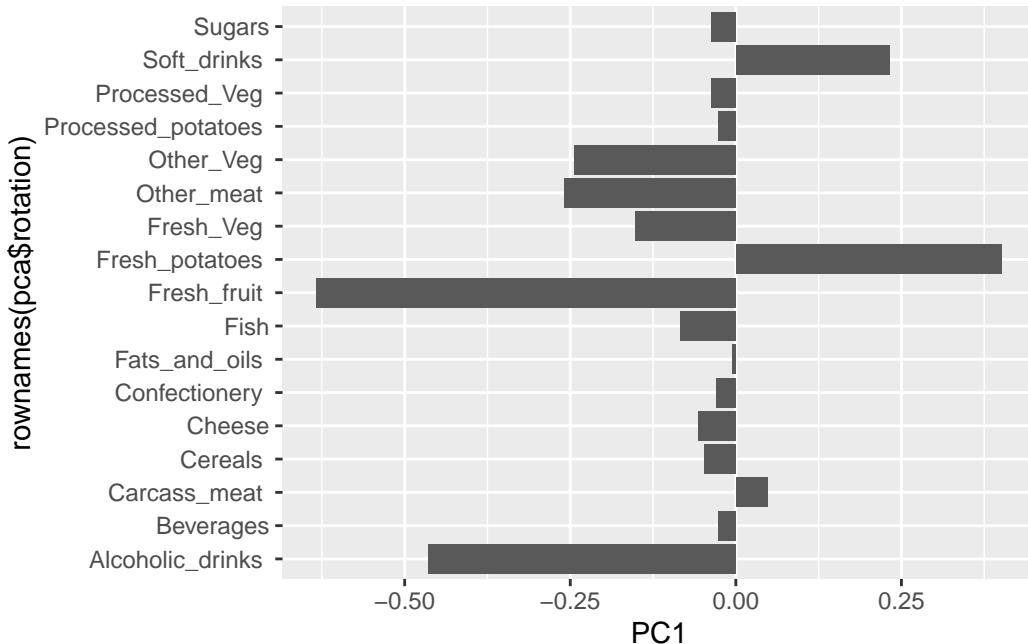


Finding out the specific foods that are contributes to the PC axis.

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001890737

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```



PCA can be a useful tool to visualize differences in large data sets.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N. Ireland consumes more potatos, soft drinks and less fresh fruit, meat, alcohol than the rest of the UK.