

Class_08_lab

Tessa Sterns PID: 18482353

Table of contents

Questions	3
Principal Component Analysis	3
Interpeting PCA Results	5
Communicating PCA Results	8
Hierachal Clustering	9
Using different methods	11
Combining methods	12
Sensitivity and Specificity	15
Prediction	15

##Background

The goal today is to use unsupervised learning and PCA to analize real world data. We will be using data on breast cancer collected through Fine Needle Aspiration. Using these techniques we hope to find trends in the data that pertain to how cell morphology is related to cancer prognosis.

##Data Import

```
wisc.df<- read.csv("WisconsinCancer.csv", row.names=1)  
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
84358402	0.1809		0.05883	0.7572	0.7813
843786	0.2087		0.07613	0.3345	0.8902
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
84348301	0.05963		0.009208	14.91	26.50
84358402	0.01756		0.005115	22.54	16.67
843786	0.02165		0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	
843786	103.40	741.6	0.1791	0.5249	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
84348301	0.6869	0.2575	0.6638		
84358402	0.4000	0.1625	0.2364		
843786	0.5355	0.1741	0.3985		
	fractal_dimension_worst				

842302	0.11890
842517	0.08902
84300903	0.08758
84348301	0.17300
84358402	0.07678
843786	0.12440

```
wisc.data <- wisc.df[,-1]

diagnosis <- factor(wisc.df$diagnosis)
```

Questions

Q1. How many observations are in this dataset?

There are `nrow(wisc.data)` observations in this data set.

Q2. How many of the observations have a malignant diagnosis?

There are `table(diagnosis)` tumors in this data set.

Q3. How many variables/features in the data are suffixed with `_mean`?

There are `length(grep("_mean", colnames(wisc.data)))` variables that have the `_mean` suffix.

Principal Component Analysis

We want to set ‘scale = T’ because it scales the pca so that those columns with high variance do not dominate.

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se

```

2.866059e+00      4.033708e+01      7.040979e-03
compactness_se     concavity_se       concave.points_se
2.547814e-02     3.189372e-02      1.179614e-02
symmetry_se        fractal_dimension_se radius_worst
2.054230e-02     3.794904e-03      1.626919e+01
texture_worst      perimeter_worst    area_worst
2.567722e+01     1.072612e+02      8.805831e+02
smoothness_worst   compactness_worst  concavity_worst
1.323686e-01     2.542650e-01      2.721885e-01
concave.points_worst symmetry_worst   fractal_dimension_worst
1.146062e-01     2.900756e-01      8.394582e-02

```

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

```
wisc.pr <- prcomp(wisc.data, scale = TRUE)
summary(wisc.pr)
```

Importance of components:

PC1	PC2	PC3	PC4	PC5	PC6	PC7
-----	-----	-----	-----	-----	-----	-----

Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

PC1 captures 44.27% of the original variance.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

There needs to be 3 principal components to capture at least 70% of the original variance.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

There needs to be 7 principle components to capture 90% of the original variance.

Interpeting PCA Results

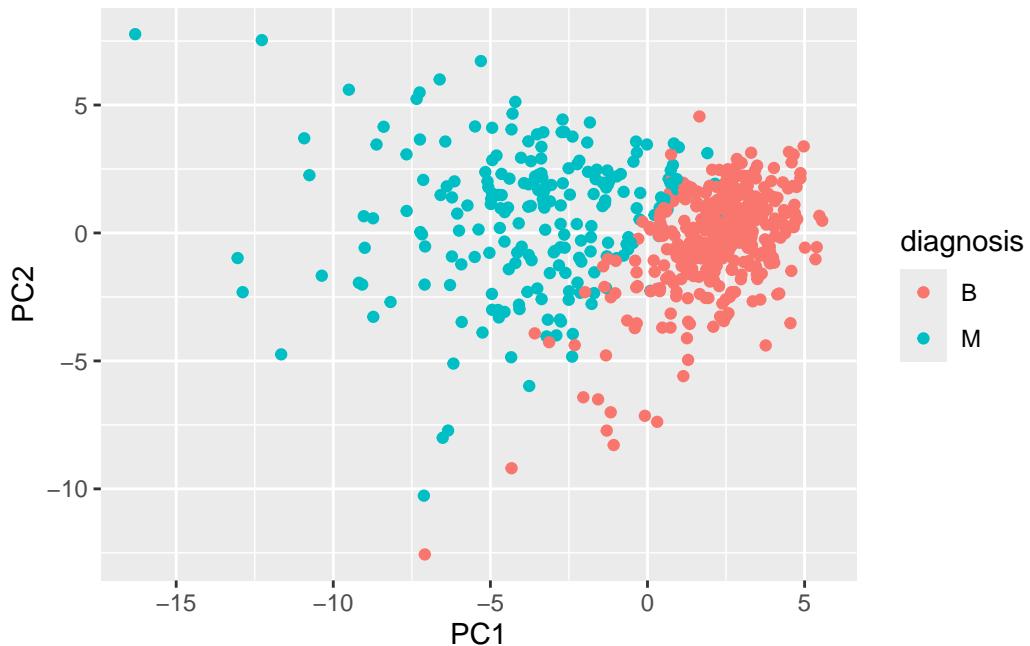
The main PC figure is called a “score plot”.

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This question is in reference to a `biplot()` which we did not make in class. However looking at the sample biplot including in the lab write up the major drawback is that they are very cluttered and difficult to interpret.

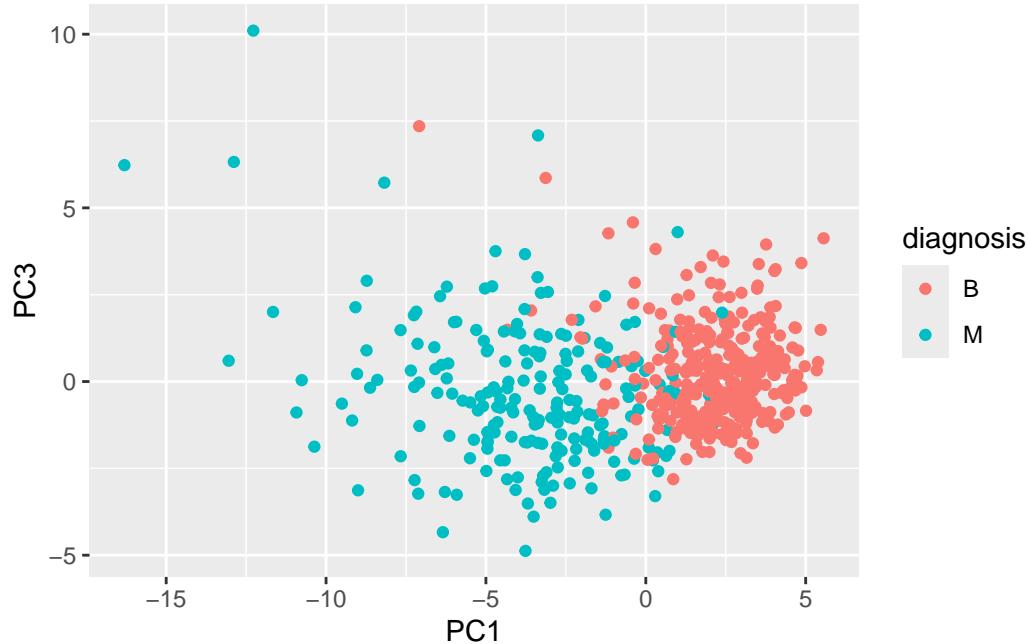
```
library(ggplot2)

ggplot(wisc.pr$x) +
  aes(x = PC1, y = PC2, col = diagnosis) +
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x) +
  aes(x = PC1, y = PC3, col = diagnosis) +
  geom_point()
```



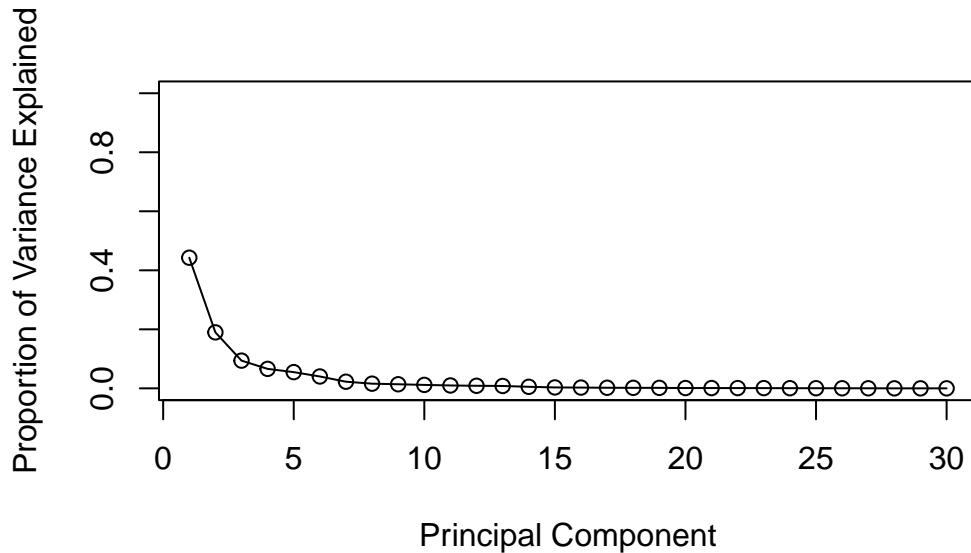
Both of these plots are very similar, the second one with PC3 mostly looks like the y-axis is shifted down. They both show that PC1 shows the most variance in the data with benign tumors appearing on the positive side of the x-axis and malignant tumors showing on the negative side of the x-axis.

```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



Communicating PCA Results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

The loading vector is `wisc.pr$rotation["concave.points_mean", "PC1"]`.

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

There needs to be five PC components to explain at least 80% of the variance.

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523

Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Hierachal Clustering

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)

data.dist <- dist(data.scaled)

wisc.hclust <- hclust(data.dist, method = "complete")
wisc.hclust
```

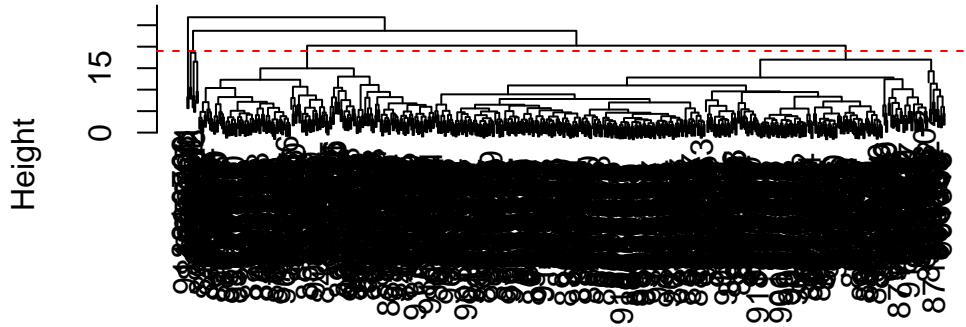
Call:
`hclust(d = data.dist, method = "complete")`

Cluster method : complete
Distance : euclidean
Number of objects: 569

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h =19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

```
wisc.hclust.clusters<- cutree(wisc.hclust, k=4)  
  
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

No the data is still ugly.

```
wisc.hclust.clusters2<- cutree(wisc.hclust, k=2)  
  
table(wisc.hclust.clusters2, diagnosis)
```

wisc.hclust.clusters2	B	M

```
1 357 210  
2   0   2
```

```
wisc.hclust.clusters10<- cutree(wisc.hclust, k=10)  
  
table(wisc.hclust.clusters10, diagnosis)
```

```
diagnosis  
wisc.hclust.clusters10   B   M  
1       12  86  
2        0  59  
3        0   3  
4      331  39  
5        0  20  
6        2   0  
7       12  0  
8        0   2  
9        0   2  
10      0   1
```

Using different methods

```
single.clus <- hclust(data.dist, method = "single")  
single.clus.cut <- cutree(single.clus, k = 2)  
table(single.clus.cut, diagnosis)
```

```
diagnosis  
single.clus.cut   B   M  
1 357 210  
2   0   2
```

```
# Not a big fan as this method as it produced a single cluster that contained both benign and malignant tumors.  
  
av.clus <- hclust(data.dist, method = "average")  
av.clus.cut <- cutree(av.clus, k = 2)  
table(av.clus.cut, diagnosis)
```

```
diagnosis  
av.clus.cut   B   M
```

```

1 357 209
2   0   3

# This method also seemed to produce a single cluster that contained both types of tumors.

wD2.clus <- hclust(data.dist, method = "ward.D2")
wD2.clus.cut <- cutree(wD2.clus, k = 2)
table(wD2.clus.cut, diagnosis)

      diagnosis
wD2.clus.cut   B   M
      1 20 164
      2 337 48

# made two distinct groups though some false positive and negatives still occurred.

```

Q13. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

Without combining PCA with clustering I think that the ward.D2 is my favorite method. Even from a pure math standpoint I think that limiting variance within a given group is a good way to determine group membership. The other methods produced a lot of dirty data that was hard to determine biological significance from.

Combining methods

Since first 3 PCA's contain most of the variance we can cluster off of them.

Behold a better tree

```

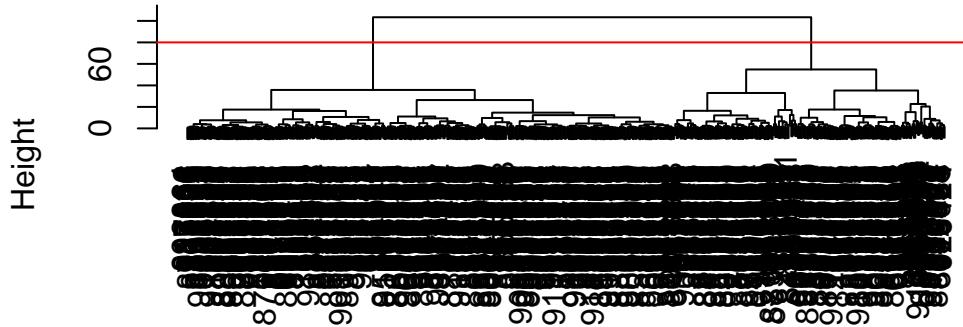
dist.pc <- dist(wisc.pr$x[,1:3])

wisc.pr.hclust<- hclust(dist.pc, method = "ward.D2")

plot(wisc.pr.hclust)
abline(h = 80, col="red")

```

Cluster Dendrogram



```
dist.pc  
hclust (*, "ward.D2")
```

To arrive at the clustering membership vector we can cut the tree at a desired height.

```
grps <- cutree(wisc.pr.hclust, k =2)  
  
table(grps)
```

```
grps  
 1 2  
203 366
```

Comparing the clustering groups to expert diagnosis.

```
table(grps, diagnosis)
```

```
diagnosis  
grps   B   M  
1    24 179  
2   333  33
```

```
## Use the distance along the first 7 PCs for clustering i.e. wisc.pr$x[, 1:7]
d2 <- dist(wisc.pr$x[,1:7])
wisc.pr.hclust2 <- hclust(d2, method="ward.D2")

wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust2, k=2)
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

Overall this method does fairly well, as good as performing clusters on PCA first.

```
table(wisc.pr.hclust.clusters, diagnosis)
```

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

```
table(wisc.pr.hclust.clusters, diagnosis)
```

wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

Using the PCA then clustering method reduced the number of false negatives we had while also increasing the number of false positives. While they both yielded fairly similar results overall using PCA had cleaner data in the end.

Sensitivity and Specificity

Sensitivity: $TP/(TP + FN)$ Specificity: $TN/(TN + FN)$

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity?

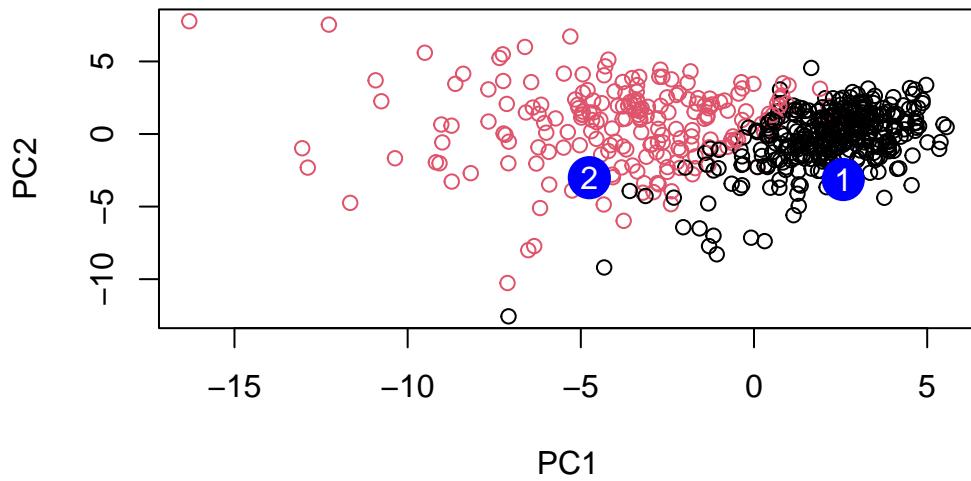
Because performing PCA before clustering resulted in fewer false negatives the PCA followed by a clustering using ward.D2 was both specific and sensitive.

Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
PC8      PC9      PC10     PC11     PC12     PC13     PC14
[1,] -0.2307350 0.1029569 -0.9272861  0.3411457  0.375921  0.1610764  1.187882
[2,] -0.3307423 0.5281896 -0.4855301  0.7173233 -1.185917  0.5893856  0.303029
PC15     PC16     PC17     PC18     PC19     PC20
[1,]  0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,]  0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
PC21     PC22     PC23     PC24     PC25     PC26
[1,]  0.1228233 0.09358453  0.08347651  0.1223396  0.02124121  0.078884581
[2,] -0.1224776 0.01732146  0.06316631 -0.2338618 -0.20755948 -0.009833238
PC27     PC28     PC29     PC30
[1,]  0.220199544 -0.02946023 -0.015620933  0.005269029
[2,] -0.001134152  0.09638361  0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col= diagnosis)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

I would prioritize patient 2, their results align with the cluster of malignant tumors whereas patient 1 is more likely benign.